

COLLABORATIVE FILTERING FOR STUDENT GRADE ANALYSIS

by

Mufan Li

A thesis submitted in conformity with the requirements
for the degree of Masters of Science
Graduate Department of Statistical Sciences
University of Toronto

© Copyright 2016 by Mufan Li

Contents

1	Introduction	1
2	Supervised Learning	2
2.1	Feed-forward Neural Networks	2
2.2	Common Techniques to Improve Training	5
2.2.1	Mini-Batches	5
2.2.2	Momentum and Adam	6
2.2.3	Dropout	7
2.2.4	Batch Normalization	9
2.3	MNIST Hand-Written Digits Example	10
3	Unsupervised Learning	13
3.1	Restricted Boltzmann Machines	13
3.2	Denoising Autoencoders	15
4	Experiment Results	18
4.1	Classification of Student Majors	18
4.2	Prediction of Course Selection	20
4.3	Prediction of Student Grades	21

List of Tables

2.1	MNIST hand-written digits test results with different neural networks.	12
4.1	Prediction of student majors test results with different neural networks.	19
4.2	Prediction of student course selection test results with different neural networks.	20
4.3	Prediction of student course grade test results with different neural networks. .	21

List of Figures

2.1	A generalized linear model represented in graphical form. In a neural network, this is also referred to as a single neuron.	3
2.2	A generalized feed-forward neural network with two hidden layers. Bias parameters are not drawn for compactness, although they are present in all forward passing nodes.	4
2.3	MNIST hand-written digits modeled using a 8 hidden layer neural network, with negative log-likelihood and classification error rate computed after each epoch.	11
3.1	A restricted Boltzmann machine (RBM) with 4 courses and 5 hidden nodes for a specific student.	14
3.2	An one layer autoencoder with 4 input nodes and 3 representation nodes. . . .	16
3.3	An one layer denoising autoencoder with 4 input nodes, 3 representation nodes, and symmetric weights. In this case, we have the third input corrupted by setting to zero.	17
4.1	Training progress for predicting student majors. Here the hyper-parameters include a learning rate of $\eta = 10^{-2}$, mini-batch size of 1000, a neural network structure of $3598 - 500 - 500 - 47$, with ReLU activation in the hidden layers, softmax activation in the output layer, negative log-likelihood objective, a dropout rate of 0.5, and trained for 200 epochs using Adam algorithm.	19

Chapter 1

Introduction

Recent developments in machine learning have made significant contributions to a wide range of fields that are not traditionally considered data science. Notably the Netflix competition have attracted a collective effort in developing models that greatly improved prediction of movie ratings by different users, creating the best movie recommendation system at the time [5]. Similar to the Netflix rating data, student grades in different courses follow the same structure, allowing the application of the same machine learning techniques. In this research course, we intend to explore several of the machine learning techniques in applications to education.

Specifically, this research project aims to apply machine learning to analyze the student grade dataset from [1], which contains complete transcripts of undergraduate students from a major Canadian University. Similar to predicting user ratings, we are able to predict the grades for courses. At the same time, we also will predict student's course and major selection choices. From the predictions, this project intends to analyze the effect of choosing easier courses on student grades, specifically by comparing the predicted grades of courses students did not take against the courses taken within the same program. By analyzing the variation in course difficulty, these results could potentially improve curriculum design for educational institutions and admission procedure for graduate programs.

This research report will be organized as follow. Chapter 2 will introduce feed-forward neural networks in detail, including several techniques that can significantly improve training results, as well as an example case using hand-written digits. Chapter 3 will introduce two unsupervised learning algorithms that can be used to for a different purpose. Chapter 4 will describe the experiments and the results we achieved using the student grade dataset.

Chapter 2

Supervised Learning

2.1 Feed-forward Neural Networks

We first consider a class of machine learning algorithms called supervised learning. In this case we have a dataset $\mathcal{D} = \{\mathbf{x}^{[n]}, \mathbf{y}^{[n]}\}, \mathbf{x}^{[n]} \in \mathbb{R}^{N_{in}}, \mathbf{y}^{[n]} \in \mathbb{R}^{N_{out}}, n \in \mathbb{N}$, with \mathbf{x} as the input, and \mathbf{y} as the label or output. We want to find a model $f(\mathbf{x}, \mathbf{w})$ such that it is the “closest” to \mathbf{y} , with \mathbf{w} the parameters in the model. This section will introduce neural networks as the model f that predicts the labels \mathbf{y} .

In the simplest case, neural networks can be reduced to a generalized linear model (GLM), where the prediction is a linear combination of inputs but passed through a non-linear function:

$$f(\mathbf{x}, \mathbf{w}) = g\left(\sum_{i=1}^{N_{in}} w_i x_i + w_0\right)$$

Here $g(\cdot)$ is a non-linear function, with \mathbf{x} is an N dimensional input vector, and \mathbf{w} is the weight vector, which includes w_0 as the bias. Common choices of $g(\cdot)$, also known as activation functions, for neural networks include the logistic (sigmoid) function, the hyperbolic tangent function, and the rectified linear unit (ReLU):

$$\begin{aligned} g_{\text{logistic}}(z) &= \frac{1}{1 + e^{-z}} \\ g_{\text{tanh}}(z) &= \tanh(z) = \frac{1 - e^{-2z}}{1 + e^{-2z}} \\ g_{\text{ReLU}}(z) &= \max(z, 0) \end{aligned}$$

where $z = \left(\sum_{i=1}^N w_i x_i + w_0 \right)$ denotes the linear combination for GLMs. Notice all of these functions have simple derivatives, and specifically logistic and hyperbolic tangent functions are monotonic and bounded by their horizontal asymptotes at infinities, which makes them great choices for binary classification problems.

Graphically, this can be represented by a series of input nodes $\{x_i\}$ connected to an output node f , with weights $\{w_i\}$ on the connections. Note bias is omitted from the graph but remains a parameter.




Figure 2.1: A generalized linear model represented in graphical form. In a neural network, this is also referred to as a single neuron.

A general feed-forward neural network is defined by recursive GLMs with different weights. For example, a neural network with two hidden layers (three layers of recursion) is defined as:

$$\begin{aligned}
 h_j^{(1)} &= g^{(1)} \left(\sum_{i=1}^{N^{(1)}} w_{ij}^{(1)} x_i + w_{0j}^{(1)} \right) \\
 h_k^{(2)} &= g^{(2)} \left(\sum_{j=1}^{N^{(2)}} w_{jk}^{(2)} h_j^{(1)} + w_{0k}^{(2)} \right) \\
 f_l &= g^{(3)} \left(\sum_{k=1}^{N^{(3)}} w_{kl}^{(3)} h_k^{(2)} + w_{0l}^{(3)} \right)
 \end{aligned}$$

where $g(\cdot)^{(\alpha)}$ is some activation function, $h_j^{(\alpha)}$ denotes the j^{th} node of the α^{th} hidden layer, $w_{ij}^{(\alpha)}$ denotes the weight for the connection of the i^{th} node of the α^{th} layer to the j^{th} node of the $(\alpha + 1)^{\text{th}}$ layer, and $N^{(\alpha)}$ denotes the number of nodes in the α^{th} layer. Additionally, let $N^{(4)}$ be the number of output nodes f_l , and $\mathbf{w} = [\mathbf{w}^{(1)} \mathbf{w}^{(2)} \mathbf{w}^{(3)}]$. Here we also note that $N^{(1)}$ is the

number of input nodes.

Graphically, this structure has a very clear representation:




Figure 2.2: A generalized feed-forward neural network with two hidden layers. Bias parameters are not drawn for compactness, although they are present in all forward passing nodes.

While most GLMs do not admit a closed-form solution, a satisfactory optimization can be achieved by the gradient descent method. In the neural network case, the optimization becomes more difficult as the number of parameters increase with the number of nodes and layers. However, we can still apply the gradient descent method and find a local optimum for the simpler neural networks. [2]

Once again we have a dataset $\mathcal{D} = \{\mathbf{x}^{[n]}, \mathbf{y}^{[n]}\}, n \in \mathbb{N}$, and we want to find a model $f(\mathbf{x}, \mathbf{w})$ such that it is the “closest” to \mathbf{y} . If the error function $E(f, \mathbf{y})$ and the model $f(\mathbf{x}, \mathbf{w})$ are differentiable with respect to \mathbf{w} , the model can be optimized by gradient descent. In other words, for any randomly initialized \mathbf{w}^0 , an improvement \mathbf{w}^{k+1} can be obtained by making a

small modification in the direction of the gradient with respect to \mathbf{w}^k :

$$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \nabla_{\mathbf{w}^k} E(f(\mathbf{x}, \mathbf{w}^k), \mathbf{y})$$

where $\eta > 0$ is hyper-parameter controlling the change of each optimization iteration, commonly called the learning rate. Note η is not part of the final model $f(\mathbf{x}, \mathbf{w})$, but it will significantly influence optimization.

In the two hidden layer neural network previously, a derivative with respect to any weight $w_{ij}^{(\alpha)}$ can be found by applying the chain rule to the derivatives. For example the derivative with respect to $w_{jk}^{(2)}$ where $j \neq 0$:

$$\begin{aligned} \text{let } z_j^{(\alpha)} &= \sum_{i=1}^{N^{(\alpha)}} w_{ij}^{(\alpha)} h_i^{(\alpha-1)} + w_{0j}^{(\alpha)} \\ \text{then } \frac{\partial E}{\partial w_{jk}^{(2)}} &= \sum_{l=1}^{N^{(4)}} \frac{\partial E}{\partial f_l} \frac{\partial f_l}{\partial z_l^{(3)}} \frac{\partial z_l^{(3)}}{\partial h_k^{(2)}} \frac{\partial h_k^{(2)}}{\partial z_k^{(2)}} \frac{\partial z_k^{(2)}}{\partial w_{jk}^{(2)}} \\ &= \sum_{l=1}^{N^{(4)}} \frac{\partial E}{\partial f_l} \frac{\partial g^{(3)}(z_l^{(3)})}{\partial z_l^{(3)}} w_{kl}^{(3)} \frac{\partial g^{(2)}(z_k^{(2)})}{\partial z_k^{(2)}} h_j^{(1)}. \end{aligned}$$

Recall $g^{(\alpha)}(\cdot)$ is selected to have a simple derivative, making the complex appearing gradient term above easy to compute.

2.2 Common Techniques to Improve Training

While the setup described in the previous section remains valid and works for simple cases, several simple techniques can significantly improve the speed and quality of optimizing the parameters in the neural network. It is important to note these techniques are also applicable in unsupervised learning techniques in Section 3.

2.2.1 Mini-Batches

Since the computational complexity of the gradient $\nabla_{\mathbf{w}^k} E$ scales linearly with the data size, and the data tends to be highly similar, it is acceptable to approximate the gradient using a small portion of the data. The resulting algorithm is to first divide up the input data into smaller batches, and then perform a gradient update for every mini-batch at random order.

The algorithm is summarized as follows:

Initialize $\mathbf{w}^0, k = 0$;

Partition dataset \mathcal{D} into N mini-batches $\{\mathcal{D}_n\}_{n=1}^N$;

repeat

 Randomize the order of mini-batches $\{\mathcal{D}_{n'}\}_{n'=1}^N$;

for $n' = 1, \dots, N$ **do**

 Use $\mathbf{x}_{n'}, \mathbf{y}_{n'} \in \mathcal{D}_{n'}$;

$\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \nabla_{\mathbf{w}^k} E(f(\mathbf{x}_{n'}, \mathbf{w}^k), \mathbf{y}_{n'})$;

$k = k + 1$;

end

until *convergence*;

Algorithm 1: The Mini-Batch Gradient Descent

Using mini-batches takes away the need to perform updates after iterating through the entire dataset, which significantly reduces the computational time.

2.2.2 Momentum and Adam

Instead of letting the gradient dictate the change in \mathbf{w}^k , the idea of momentum is to let the gradient dictate the rate of change. If \mathbf{w}^k is interpreted as a coordinate, and each optimization iteration as velocity, momentum can be seen as using the gradient as acceleration instead of velocity. This allows the optimization to accumulate speed in a consistent direction of the gradient, while making it harder to slow down and converge to a poor local minimum.

At the same time, momentum also improves the speed of convergence. It is well known that for a convex optimization problem, the gradient descent method will converge at a rate of $\mathcal{O}(k^{-1})$, where k is the number of iterations. The momentum technique is a special case of Nesterov Acceleration [12], where the rate of convergence is $\mathcal{O}(k^{-2})$ for a strongly convex objective function and a Lipschitz continuous gradient. Since the neural network objective function is highly non-convex, the momentum method tends to perform better in practice.

The formulation starts with a velocity vector \mathbf{v}^0 initialized to zero, and the rest is similar:

$$\begin{aligned}\mathbf{v}^{k+1} &= \theta \mathbf{v}^k - \eta \nabla_{\mathbf{w}^k} E(f(\mathbf{x}, \mathbf{w}^k), \mathbf{y}) \\ \mathbf{w}^{k+1} &= \mathbf{w}^k + \mathbf{v}^{k+1}\end{aligned}$$

where $\theta \in [0, 1]$ is the hyper-parameter deciding the preservation of momentum. Here choosing

a larger θ would result in a stronger preservation of the velocity vector \mathbf{v} , which then retains more momentum.

Recently, many new algorithms are developed for adaptively adjusting the gradient and momentum. Here we highlight one specific algorithm called Adam [8], where the authors here choose to work with an exponential moving average of the gradient and its second moment. In this case, a momentum-like result is achieved with the moving average, and the direction of the gradient is scaled by the curvature as well.

To specific the algorithm we require a learning rate η , exponential moving average decay parameters β_1, β_2 , and lower bound on second moment ϵ . Here we denote $(\mathbf{g}^k)^2 = \mathbf{g}^k \odot \mathbf{g}^k$ as the element wise square, and $\beta_1^k = (\beta_1)^k$ as the k^{th} power. The Adam algorithm then follows

```

Initialize  $\mathbf{w}^0, k = 0, \mathbf{m}^0 = 0, \mathbf{v}^0 = 0$ ;
repeat
    Use  $\mathbf{x}, \mathbf{y} \in \mathcal{D}$  ;
    Compute gradient
     $\mathbf{g}^{k+1} = \nabla_{\mathbf{w}^k} E(f(\mathbf{x}, \mathbf{w}^k), \mathbf{y})$  ;
    Compute the first and second moment exponential moving average
     $\mathbf{m}^{k+1} = \beta_1 \mathbf{m}^k + (1 - \beta_1) \mathbf{g}^{k+1}$  ;
     $\mathbf{v}^{k+1} = \beta_2 \mathbf{v}^k + (1 - \beta_2) (\mathbf{g}^{k+1})^2$  ;
    Correct for bias
     $\hat{\mathbf{m}}^{k+1} = \mathbf{m}^{k+1} / (1 - \beta_1^k)$  ;
     $\hat{\mathbf{v}}^{k+1} = \mathbf{v}^{k+1} / (1 - \beta_2^k)$  ;
    Update parameters
     $\mathbf{w}^{k+1} = \mathbf{w}^k - \eta \hat{\mathbf{m}}^{k+1} / (\sqrt{\hat{\mathbf{v}}^{k+1}} + \epsilon)$  ;
     $k = k + 1$  ;
until convergence;
```

Algorithm 2: Adam Algorithm

Observe when \mathbf{g}^k is consistently small, we would have a small \succ^k , however it will be divided by a small \asymp^k as well, adjusting for slow down in convergence.

2.2.3 Dropout

Another common issue for neural networks is over-fitting. Due to the large number of parameters, a neural network can tend to “store” the entire dataset into its parameters, hence

over-fitting the training dataset. While regularization and early stopping are used for preventing this issue, we explore a simple yet highly effective technique called dropout [16].

The motivation for dropout comes from ensemble methods, where multiple model predictions are averaged, with each model weighted by its posterior probability given the data. Ensembles are highly successful when a large number of distinct models can be generated with relative low computation. A notable example is random forests [3] where each model is a simple decision tree. For neural networks, the simplest method to create distinct models is by considering different model architectures. However this is difficult to apply directly due to the computational cost of optimizing even one neural network. More recently, authors in [6] showed that dropout objective is minimizing the Kullback-Leibler divergence between an approximate model and a deep Gaussian process model in [4]. Mathematically, this explains how dropout can avoid model over-fitting.

Dropout is the method that attempts to incorporate random neural network architecture into the same training procedure. The technique specifically refers to “dropping out” some of the hidden units with some probability $p \in [0, 1]$, creating a random structure for each gradient descent iteration.

Specifically recall the feed-forward neural network, where we originally had

$$h_k^{(\alpha)} = g^{(\alpha)} \left(\sum_{j=1}^{N^{(\alpha)}} w_{jk}^{(\alpha)} h_j^{(\alpha-1)} + w_{0k}^{(\alpha)} \right)$$

Here for each gradient descent iteration, we introduce a list of Bernoulli random variables $r_j^{(\alpha-1)} \sim \text{Bernoulli}(1 - p)$, and modify the previous equation to

$$h_k^{(\alpha)} = g^{(\alpha)} \left(\sum_{j=1}^{N^{(\alpha)}} w_{jk}^{(\alpha)} h_j^{(\alpha-1)} r_j^{(\alpha-1)} + w_{0k}^{(\alpha)} \right)$$

Observe that with probability p , each hidden node $h_j^{(\alpha-1)}$ will be set to zero, hence creating a random structure. We also observe that since if the node $h_j^{(\alpha-1)}$ is dropped, we have that the parameter $w_{jk}^{(\alpha)}$ stays unchanged for this iteration, since the hidden node $h_j^{(\alpha-1)}$ does not affect the objective function for this iteration. Similarly, all the parameters $w_{ij}^{(\alpha-1)}, \forall i$ intended for the nodes leading up to $h_j^{(\alpha-1)}$ remain unchanged as well.

2.2.4 Batch Normalization

Finally, a significant hurdle in training neural networks is the constantly changing distribution of each layer's input during training, as the parameters of the previous layer changes. This results in the gradient descent method struggling to optimize each layer at the same rate. Previously, most optimization is done by manually tuning the learning rate separately for each layer, which is notoriously difficult. Authors in [7] propose a simple adjustment to resolve the issue, namely batch normalization. The idea is to normalize each layer's input to have a standard mean and variance before activation, controlling the distribution of different layers.

Consider our feed-forward activation as before

$$z_k^{(\alpha)} = \sum_{j=1}^{N^{(\alpha)}} w_{jk}^{(\alpha-1)} h_j^{(\alpha-1)} r_j^{(\alpha-1)} + w_{0k}^{(\alpha)}$$

$$h_k^{(\alpha)} = g^{(\alpha)} \left(z_k^{(\alpha)} \right)$$

Now instead of using $z_k^{(\alpha)}$ directly, we want to normalize it to $\hat{z}_k^{(\alpha)}$ where $\mathbb{E}\hat{z}_k^{(\alpha)} = 0$ and $Var[\hat{z}_k^{(\alpha)}] = 1$. To approximate this in a mini-batch setting, we consider all the values before activation in the same mini-batch $\{z_k^{(\alpha)}(i)\}_{i=1}^B$, where B is the size of the mini-batch, and we normalize it according to its sample mean and variance. Specifically

Before each activation

Compute the pre-activation values for each i in the mini-batch

$$z_k^{(\alpha)}(i) = \sum_{j=1}^{N^{(\alpha)}} w_{jk}^{(\alpha-1)} h_j^{(\alpha-1)}(i) r_j^{(\alpha-1)} + w_{0k}^{(\alpha)};$$

Compute the sample mean and variance of the batch

$$\mu_k^{(\alpha)} = \frac{1}{B} \sum_{i=1}^B z_k^{(\alpha)}(i);$$

$$\sigma_k^{(\alpha)} = \sqrt{\frac{1}{B} \sum_{i=1}^B \left[z_k^{(\alpha)}(i) - \mu_k^{(\alpha)} \right]^2};$$

Normalize the pre-activation values and compute activation outputs

$$\hat{z}_k^{(\alpha)}(i) = \left[z_k^{(\alpha)}(i) - \mu_k^{(\alpha)} \right] / \sigma_k^{(\alpha)};$$

$$h_k^{(\alpha)} = g^{(\alpha)} \left(\hat{z}_k^{(\alpha)} \right)$$

Algorithm 3: The Batch Normalization Algorithm

2.3 MNIST Hand-Written Digits Example

The Mixed National Institute of Standards and Technology (MNIST) dataset [10] is a collection of images of hand-written digits from various sources, with each image labeled the correct digit. The dataset contains 60,000 images for training (fitting), and 10,000 images for testing. The images are 28x28 in resolution, hence making $N^{(1)} = 784$ dimensions in input.

The data labels are changed to use the 1-of-K encoding scheme, where the label \mathbf{y} is a binary vector of size K , with only one element taking a value of one. In this case, given 10 possible digits, we have a vector of size $N_{out} = 10$, where we coincidentally also chose 10 layers. For example, a possible scheme can label the digit “3” using the vector $[0, 0, 1, 0, \dots]$ where only the 3rd index is a “1”.

To best model this type of label vector, the softmax function is chosen for the output layer:

$$f_l = g^{(9)}(z_l^{(9)}) = \frac{\exp(z_l^{(9)})}{\sum_{k=1}^{N^{(10)}} \exp(z_k^{(9)})}$$

where $z_l^{(9)} = \sum_{k=1}^{N^{(9)}} w_{kl}^{(9)} h_k^{(8)} + w_{0l}^{(9)}$ is a linear combination of the final hidden layer. Since the denominator normalizes the sum, the f_l now adds up to one, and a “perfect” output is exactly the 1-of-K encoded label. If f_l is modeled as the probability of the image being digit l , suppose the correct digit is m , then the likelihood of making the correct prediction is:

$$L(\mathbf{f}, \mathbf{y}) = f_m = \prod_{l=1}^{N^{(10)}} f_l^{y_l}$$

since $y_m = 1$ is the only non-zero term in the label vector. We can then define the error function as negative log-likelihood:

$$E(\mathbf{f}, \mathbf{y}) = -\log \prod_{l=1}^{N_{out}} f_l^{y_l} = -\sum_{l=1}^{N_{out}} y_l \log f_l$$

where N_{out} is the number of output nodes, and minimizing E is equivalent to maximizing likelihood. Note taking the logarithm creates an error function with much simpler derivative, hence simplifying the gradient descent method.

In the following experiment, several different architectures are used to model the MNIST digits. For example, we denote $N^{(2)} = N^{(3)} = \dots = N^{(9)} = 500$ nodes in the hidden layers, creating a structure of $784 - 500 - \dots - 500 - 10$ ($N^{(1)} - N^{(2)} - \dots - N^{(10)}$) nodes in each

layer. We also chose $g^{(1)}(\cdot) = \dots = g^{(9)}(\cdot) = g_{\text{ReLU}}(\cdot)$ in the hidden layers, and softmax for the output layer. For all the experiments on MNIST dataset, we use mini-batches, Adam, dropout, and batch normalization altogether. The hyper parameters were chosen as $\eta = 10^{-1}$ for the learning rate, $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$ for batch normalization, and $p = 0.3$ for dropout rate. We also chose to update the weight vector \mathbf{w}^k once for every 100 samples of digits, also known as a mini-batch.

After training (optimizing) for 100 epochs, with each epoch denoting one complete run through of the training dataset, we reach a test error rate of 0.43% (Figure 2.3). We highlight

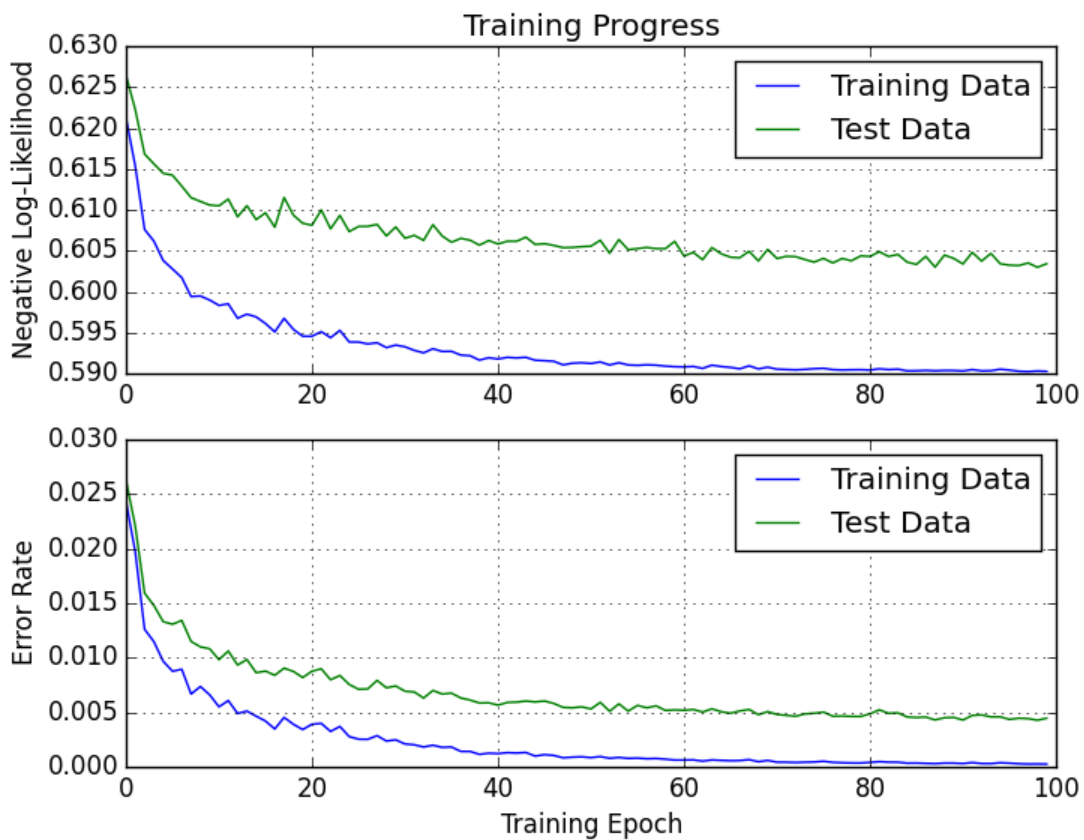


Figure 2.3: MNIST hand-written digits modeled using a 8 hidden layer neural network, with negative log-likelihood and classification error rate computed after each epoch.

that the impressive classification results on MNIST digits is highly motivating for applications on different problems.

We also note that scaling to higher number of layers and nodes yields diminishing returns. From Table 2.1, we observe that 6 hidden layers can achieve practically the same result as 8 hidden layers.

Hidden Layers	Nodes in Each Hidden Layer	Test Error Rate
4	500	0.70%
4	1000	0.57%
6	500	0.44%
8	500	0.43%

Table 2.1: MNIST hand-written digits test results with different neural networks.

While this type of neural networks is feed-forward, which mean it is limited to only supervised type problems where the data structure is consistent and a prediction target (label) is provided for each sample. For a collaborative filtering type problem, the inference is often made within the data structure itself, which makes an unsupervised learning problem. Feed-forward neural networks also fail to fully utilize the datasets that are partly labeled, known as semi-supervised problems. These problems would require other variations of neural networks with different methods for inference.

Chapter 3

Unsupervised Learning

3.1 Restricted Boltzmann Machines

On the other hand, restricted Boltzmann machines (RBM) is a completely different approach to problems without labels. RBM is a type of unsupervised learning algorithm, for there are no labels to “supervise” the learning. The purpose of unsupervised algorithms are to find structural patterns within the data itself. In this case, we are interested in the relationships between the performance in different courses, and how this helps us predict the grades.

A RBM is a Markov random field in the form of a bipartite graph, where the joint probability follows a Boltzmann type distribution. The bipartite graph structure creates two layers without internal connections. One layer, called the visible layer, contain the input data; in this case, the visible values are the grades of each student. These nodes are connected to the other layer, called the hidden layer, with symmetrical weighted connections.

Suppose the graph have N visible nodes and M hidden nodes, with each visible node denoted v_i , hidden nodes denoted h_j , weights between two nodes w_{ij} , b_i and a_j be bias parameters, and σ_i be the standard deviation of grades for each course. Here each visible node v_i represents the grade for course i , where a specific student is fixed. Let $\theta = \{w_{ij}, a_j, b_i, \sigma_i\} \forall i, j$, $\mathbf{v} = \{v_i\} \forall i$, and $\mathbf{h} = \{h_j\} \forall j$ denote the collections. Additionally, we let the hidden nodes only take on binary values, i.e. $v_i \in \mathbb{R}, h_j \in \{0, 1\}$. We can then define the energy function




Figure 3.1: A restricted Boltzmann machine (RBM) with 4 courses and 5 hidden nodes for a specific student.

and the joint distribution for the graph:

$$E(\mathbf{v}, \mathbf{h}|\theta) = \sum_{i=1}^N \frac{(b_i - v_i)^2}{2\sigma_i} - \sum_{i=1}^N \sum_{j=1}^M w_{ij} h_j \frac{v_i}{\sigma_i} - \sum_{j=1}^M a_j h_j$$

$$P(\mathbf{v}, \mathbf{h}|\theta) = \frac{\exp[-E(\mathbf{v}, \mathbf{h}|\theta)]}{\mathcal{Z}}$$

where \mathcal{Z} is the partition function normalizing the distribution. After marginalizing over the hidden nodes \mathbf{h} , we can find the gradient of the likelihood function with respect to the parameters θ to perform steepest descent optimization. Finding the gradient requires the use of Gibbs sampling, although [15] showed the approximate gradient after very few iterations of Gibbs sampling is sufficient for optimization.

$$\frac{\partial P(\mathbf{v}|\theta)}{\partial w_{ij}} = \mathbf{E}_{\text{data}}(v_i h_j) - \mathbf{E}_{\text{model}}(v_i h_j)$$

where \mathbf{E}_{data} refers to expectation of observing the case within data, and $\mathbf{E}_{\text{model}}$ is the expectation of the current model with parameters θ . Instead of using Gibbs sampling until convergence to find $\mathbf{E}_{\text{model}}$, [15] uses k iterations for a very good approximation of the gradient. This method is referred to contrastive divergence (CD) by the authors in [15], where CD- k refers to k iterations used in Gibbs sampling. As a result, we have a very good algorithm to optimize the RBM for likelihood.

To perform inference on a missing grade value, one simply includes an additional “visible”

node v_p , where the value is not known, but can be determined by the energy function:

$$\begin{aligned} P(v_p|\mathbf{v}) &\propto \sum_{\mathbf{h}} \exp[-E(v_p, \mathbf{v}, \mathbf{h})] \\ &= \prod_{j=1}^M \left(1 + \exp \left[\sum_{i=1}^N w_{ij} v_i \right] \right) \end{aligned}$$

Alternatively, we can also treat the RBM weights as the weights for an autoencoder, which will be discussed in the next section.

3.2 Denoising Autoencoders

Another approach to unsupervised learning is using autoencoders (AEs), specifically in this case we will introduce the denoising autoencoders (DAEs) in [17]. The autoencoder is a compression model of input data, such that a high dimensional input can be encoded as a low dimensional representation, where the data can be reconstructed from the representation using a decoder.

For this problem we consider a dataset $\mathcal{D} = \{\mathbf{x}^{[n]}\}, \mathbf{x}^{[n]} \in \mathbb{R}^{N_{in}}, n \in \mathbb{N}$, with only \mathbf{x} as the input. We also define a desired feature $\mathbf{h} \in \mathbb{R}^{N_{feat}}$ with $N_{feat} < N_{in}$, and an encoder-decoder pair $f(\mathbf{x}, \mathbf{w}^{(1)}), g(\mathbf{h}, \mathbf{w}^{(2)})$ such that the reconstruction $\hat{\mathbf{x}} = g \circ f(\mathbf{x}) \approx \mathbf{x}$. This results in a forced compression of input \mathbf{x} into lower dimensional \mathbf{h} , and in the process, the parameterization \mathbf{w} retains further information about the data structure.

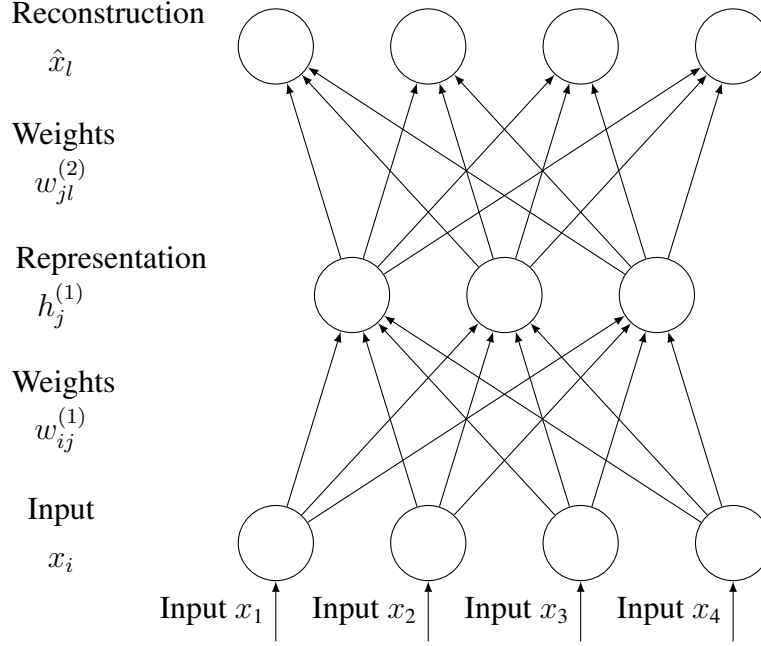


Figure 3.2: An one layer autoencoder with 4 input nodes and 3 representation nodes.

With this setup, we have a neural network described as in Figure 3.2. In this structure, we can optimize for the optimal parameters $\{\mathbf{w}^{(1)}, \mathbf{w}^{(2)}\}$ using the gradient descent approach from feedforward neural networks 2.1. In practice, tied weights condition $\mathbf{w}^{(1)} = [\mathbf{w}^{(2)}]^\top$ is often enforced to start the optimization. Observe that when the weights are tied and the non-linear activation function is sigmoid, we have a striking similarity with the RBM: the representation \mathbf{h} is exactly the probability of binary hidden layer sampled as 1.

However as [17] explained, pure compression retains insufficient information, especially when compared to RBMs; therefore the authors introduced a new optimization criterion: reconstruction from noisy inputs. Formally, we have a corruption function q that creates noisy inputs $\mathbf{v} = q(\mathbf{x})$. A popular choice of q is to randomly set a fraction of the input dimensions to zero.

To motivate denoising autoencoders, we consider an example with 2 inputs, i.e. $\mathbf{x} = \{x_1, x_2\}$, and let $x_1 \approx \phi(x_2)$ for some bijection ϕ . When x_1 is set to zero due to corruption, it remains possible to reconstruct x_1 by learning the relationship between x_1 and x_2 , which gives us $\hat{x}_1 = \phi(x_2)$. Similarly, when x_2 is corrupted, ideally we can have $\hat{x}_2 = \phi^{-1}(x_1)$.

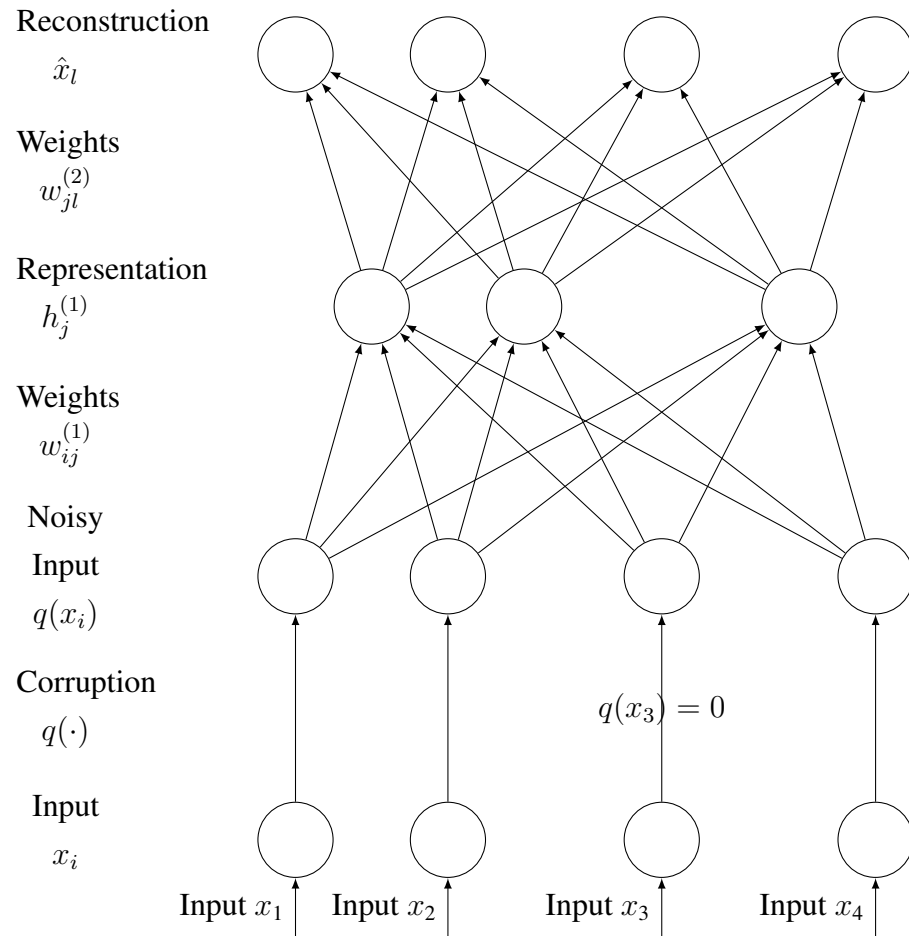


Figure 3.3: An one layer denoising autoencoder with 4 input nodes, 3 representation nodes, and symmetric weights. In this case, we have the third input corrupted by setting to zero.

Chapter 4

Experiment Results

4.1 Classification of Student Majors

Our first experiment involves the prediction of a student's choice of majors. Specifically, we define the input data to be the course grades of the student in the first two years of studies, and the label of each student to be the major. Specifically, we filter for only the courses for each student during the two years of enrollment, and predict the most department of which the most courses were taken from after the first two years. Since the task at hand is a classification problem, we will use the negative log-likelihood loss function introduced in Section 2.3.

We can see the training progress in Figure 4.1. Here we can observe the best error rate was reached around 45%. We can also observe that even with using Dropout, there remains a significant difference between training results and test results. It is highly probable that the training data is simply not rich enough to capture the structure of student grades. To confirm the suspicion, we can compare the classification results for different architecture choices.

From Table 4.1, we can observe that even with one hidden layer of 100 nodes, we have already reached very close to the best result. In fact, with hidden nodes and layers, the results can actually get worse. These results highlight the difficulty of predicting student majors only based on the first two years of courses.

Instead of using the student's grades, we can also attempt to predict with only the list of courses the student took and achieve the exact same result as the table below. This is important because it implies that the neural network did not taken into account the information with the student's grades.

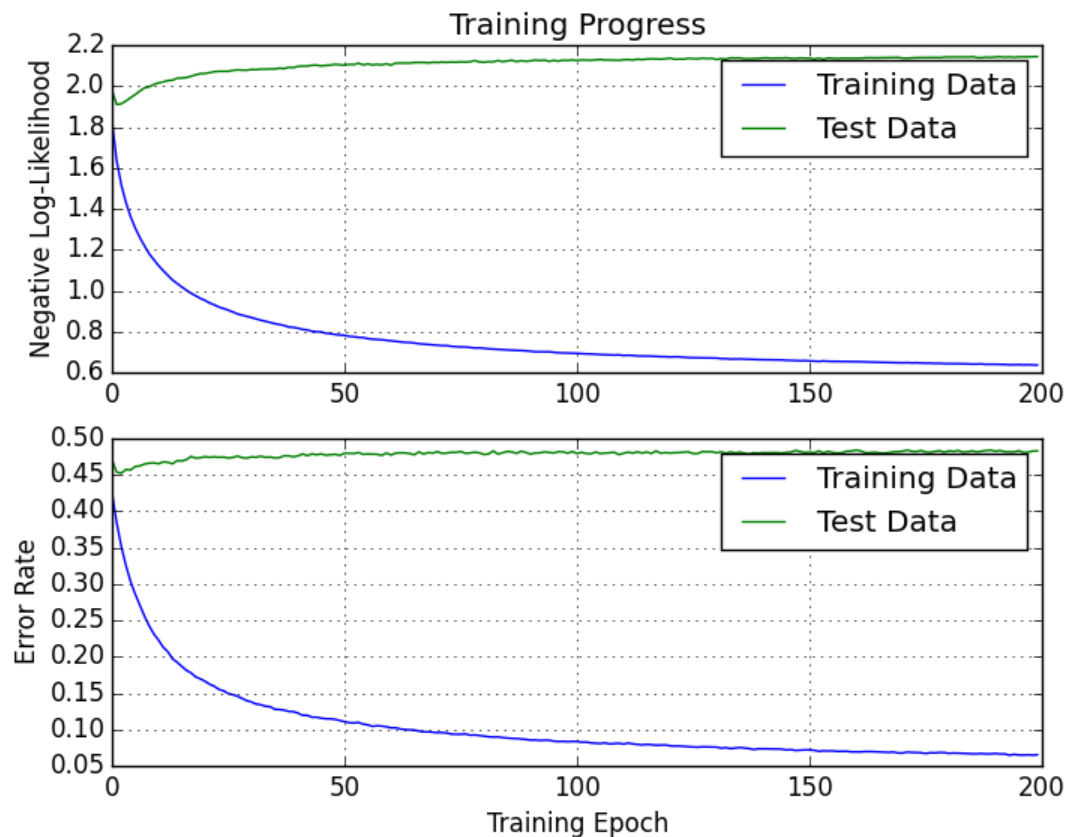


Figure 4.1: Training progress for predicting student majors. Here the hyper-parameters include a learning rate of $\eta = 10^{-2}$, mini-batch size of 1000, a neural network structure of $3598 - 500 - 500 - 47$, with ReLU activation in the hidden layers, softmax activation in the output layer, negative log-likelihood objective, a dropout rate of 0.5, and trained for 200 epochs using Adam algorithm.

Hidden Layers	Nodes in Each Hidden Layer	Test Error Rate
1	100	45.3%
1	500	45.8%
1	1000	45.8%
2	100	45.5%
2	500	44.8%
2	1000	45.1%
4	500	45.5%
6	500	47.8%
8	500	49.1%

Table 4.1: Prediction of student majors test results with different neural networks.

4.2 Prediction of Course Selection

We then move onto predicting whether or not a student will select to take an upper year course. Specifically, using only the courses during the first two years of enrollment, we want to predict all the courses that a student will take. In this case we will need to choose use a slightly different likelihood function as objective.

$$E(\mathbf{f}, \mathbf{y}) = -\frac{1}{N_{taken}} \sum_{l=1}^{N_{out}} y_l \log f_l - \frac{1}{N_{out} - N_{taken}} \sum_{l=1}^{N_{out}} (1 - y_l) \log(1 - f_l)$$

Observe in the above equation, what we have instead of the previous negative log-likelihood, we have an objective function equally punishing to both false positive and false negative errors. That is, the objective will equally punish both when the student **does not take** a course and we predict **true**, and when the student **does take** a course and we predict **false**. Evaluating such an objective will discourage the neural network to converge to a degenerate case where prediction is all true or all false.

Hidden Layers	Hidden Nodes	False Negative	False Positive
1	100	1.63%	44.86%
1	500	1.76%	45.87%
1	1000	2.05%	45.38%
2	100	2.04%	44.86%
2	500	1.76%	42.81%
2	1000	2.02%	41.64%
4	500	1.90%	33.73%
4	1000	2.07%	34.40%
6	500	2.16%	30.52%
6	1000	2.41%	29.51%

Table 4.2: Prediction of student course selection test results with different neural networks.

Observe the results in 4.2, we have that limiting the false negative rate is a much simpler task than false positives. This is a clear since a student often takes around 40 courses in an undergraduate degree, while there are over 4000 course choices to predict from. The fact that we are limiting false positives below 50% alone is implying that our model cannot be guessing to achieve a 2% false negative rate.

Unlike predicting student majors, prediction errors here clearly scales with a larger neural network. However, a false positive error rate of 30% implies that the current model tends to predict the student will take a very large pool of courses. Depending on the exact criterion of

a good prediction, it may be helpful to place regularization on the total number of courses the model predicts the student to take.

4.3 Prediction of Student Grades

Similar to predicting whether a student will take a course in the future, we are also interested in predicting the grade of such a course. To make such a prediction, we opt for unsupervised learning algorithms. We have introduced restricted Boltzmann machine (RBM) and denoising auto-encoder (DAE) in Section 3, and these methods fit the task of filling in missing data.

To frame the problem of predicting a student’s grade in the form of reconstruction, we simply assume the course that we are trying predict is a “missing” data point. With RBM or DAE reconstructing the missing data point, we are able to predict the missing grade.

We train both the RBM and DAE with a learning rate of $\eta = 10^{-4}$, momentum constant of $\theta = 0.9$, batch size of 100, and for 100 epochs using the root mean squared error (RMSE) objective. For the RBM, we progressively increased the number of Gibbs iterations as 2, 5, 15 over training time. For the DAE, we used a noise level of $p = 0.3$.

Hidden Nodes	RBM RMSE	DAE RMSE
20	11.52%	20.57%
50	11.77%	18.64%
100	13.29%	17.98%
200	14.81%	19.38%

Table 4.3: Prediction of student course grade test results with different neural networks.

From Table 4.3, we can see that not only does more hidden nodes not necessarily help in this case, it may made the network more difficult to train. At the same time, we can see that predicting grades is once again a difficult task, as even the best RMSE is at 11.52%. Suppose that the error is normally distributed, this implies at least 30% of the time, the prediction is worse than 11.52%. Further improvements are definitely possible to these models with other types of modifications by stacking RBMs and DAEs [13, 17], or alternatively using a variational auto-encoder [9].

Bibliography

- [1] Michael A Bailey, Jeffrey S Rosenthal, and Albert H Yoon. Grades and incentives: assessing competing grade point average measures and postgraduate outcomes. *Studies in Higher Education*, (ahead-of-print):1–15, 2014.
- [2] C Bishop. Pattern recognition and machine learning (information science and statistics), 1st edn. 2006. corr. 2nd printing edn, 2007.
- [3] Leo Breiman. Random forests. *Machine learning*, 45(1):5–32, 2001.
- [4] Andreas C Damianou and Neil D Lawrence. Deep gaussian processes. *arXiv preprint arXiv:1211.0358*, 2012.
- [5] Andrey Feuerverger, Yu He, Shashi Khatri, et al. Statistical significance of the netflix challenge. *Statistical Science*, 27(2):202–231, 2012.
- [6] Yarin Gal and Zoubin Ghahramani. Dropout as a bayesian approximation: Appendix. *arXiv preprint arXiv:1506.02157*, 2015.
- [7] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *arXiv preprint arXiv:1502.03167*, 2015.
- [8] Diederik Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [9] Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [10] Yann LeCun, Corinna Cortes, and Christopher JC Burges. The mnist database of handwritten digits, 1998.
- [11] Andriy Mnih and Ruslan Salakhutdinov. Probabilistic matrix factorization. In *Advances in neural information processing systems*, pages 1257–1264, 2007.

- [12] Yurii Nesterov. *Introductory lectures on convex optimization: A basic course*, volume 87. Springer Science & Business Media, 2013.
- [13] Ruslan Salakhutdinov. *Learning deep generative models*. PhD thesis, University of Toronto, 2009.
- [14] Ruslan Salakhutdinov and Andriy Mnih. Bayesian probabilistic matrix factorization using markov chain monte carlo. In *Proceedings of the 25th international conference on Machine learning*, pages 880–887. ACM, 2008.
- [15] Ruslan Salakhutdinov, Andriy Mnih, and Geoffrey Hinton. Restricted boltzmann machines for collaborative filtering. In *Proceedings of the 24th international conference on Machine learning*, pages 791–798. ACM, 2007.
- [16] Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *The Journal of Machine Learning Research*, 15(1):1929–1958, 2014.
- [17] Pascal Vincent, Hugo Larochelle, Isabelle Lajoie, Yoshua Bengio, and Pierre-Antoine Manzagol. Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion. *The Journal of Machine Learning Research*, 11:3371–3408, 2010.